Fine Print

Reliability and Limitations

A guide to how Fine Print is engineered for trustworthiness, and what it can and cannot do.

What Fine Print is, and isn't

Fine Print is a first-pass triage tool. It helps a non-lawyer quickly understand what a contract says, which clauses deviate from the norm, and what standard protections might be missing, so they know what to ask about before signing or before paying for a lawyer's time.

Fine Print is not a lawyer and not a substitute for legal review. It provides legal information, not legal advice. This distinction is not a disclaimer bolted on at the end; it shaped the entire product, for two reasons.

- **No ground truth.** An LLM's legal judgment cannot be verified against a validated dataset of lawyer-reviewed contracts, because no such dataset exists here. The false-negative rate, how often a genuinely predatory clause is marked standard, therefore cannot be stated. A tool that cannot quantify its own error rate must not present itself as authoritative.
- **Regulatory precedent.** The FTC's 2025 action against DoNotPay established that AI tools marketed as substitutes for professional legal services, without validation, create consumer harm. Fine Print is deliberately positioned as an informational aid, explicitly not a legal-advice product.

The most dangerous failure mode for a tool like this is false confidence: a user trusting a clean-looking analysis that silently missed something. Every mechanism below exists to surface uncertainty rather than hide it.

Reliability mechanisms

1. Verbatim quote verification

Problem. LLMs hallucinate. An analysis might attribute a quote to a contract that does not actually appear in it; a user reading that fabricated quote would believe their contract contains language it does not.

Mechanism. Every quote the model returns is checked character-by-character against the source text. Matching is normalized for case, whitespace, and smart-quote variance so genuine quotes are not falsely rejected over a curly apostrophe, but a quote that does not substantively appear in the source is flagged unverified. The interface marks those clauses and warns the reader to check against the original.

Limitation. Verification runs only for pasted text, where the source is available. For PDF and image uploads the model reads the document directly and there is no separate text to compare against. This is disclosed, not hidden.

2. Self-consistency (deep check)

Problem. A single LLM run can be unstable; the same contract might be scored differently on a different run. One run hides this instability behind one confident-looking answer.

Mechanism. When deep check is enabled, the contract is analyzed twice, once at temperature 0 and once at temperature 0.4, and the two analyses are reconciled clause by clause. Where both runs agree, confidence is boosted. Where they disagree, the more cautious rating is shown, confidence is dropped, and the clause is flagged uncertain with the alternate rating disclosed. A clause seen by only one run is flagged partial.

Why temperature 0.4 on the second pass. At temperature 0 both runs are nearly identical, making the consistency check theater. A moderate temperature lets the second pass genuinely diverge, so agreement becomes a real signal and disagreement surfaces honest uncertainty.

Why opt-in, not default. Deep check doubles cost and latency. Imposing that on every analysis, including a one-page balanced lease that does not need it, is the wrong default. Making it a user choice respects that reliability has a cost and lets it be spent where the stakes justify it.

Why the more cautious rating wins. For a consumer-protection tool, a false alarm costs the user a few minutes of attention; a miss could cost them materially. That asymmetry justifies erring toward caution.

3. Coverage and truncation detection

Problem. A long contract can exceed the model's maximum output length, cutting off the analysis mid-result and silently dropping clauses the user never learns about.

Mechanism. Truncation is detected via the API stop reason and incomplete JSON. A character-level recovery parser salvages every complete clause from a truncated response and discards partial ones, then flags the result. The interface shows an explicit banner that coverage may be partial and suggests analyzing in smaller sections.

4. Input validation and non-contract detection

Problem. Garbage in, confident garbage out. Pasting a recipe should not produce a fabricated contract analysis.

Mechanism. Text length is bounded, and the model is instructed to declare when the input is not a contract. The interface then shows a helpful message instead of inventing an analysis.

5. Deterministic, transparent scoring

Problem. If the health score came from the LLM it could vary run to run and contradict the clauses on screen.

Mechanism. The 0 to 100 score is computed client-side by a fixed formula from clause and gap counts, never by the model. The same clauses always produce the same score, and the score always matches what is displayed. Each deduction type is capped so the scale stays meaningful.

6. Retry, timeout, and graceful failure

Transient failures such as rate limits and overload retry with exponential backoff; calls time out cleanly at 55 seconds; and infrastructure errors such as payload-size and timeout codes are translated into plain-English guidance.

7. Confidence signaling

Every clause carries a confidence score; the overall analysis confidence is surfaced as a labeled badge. Low-confidence clauses, those that are ambiguous or jurisdiction-dependent, are individually flagged.

Known limitations

- **No jurisdiction awareness.** A clause legal in one jurisdiction may be void in another. Fine Print does not apply local law.
- **Accuracy is unmeasured.** Confidence scores reflect the model's self-assessment, not a validated accuracy rate.
- **Reference templates are general.** Gap detection compares against general standards, not jurisdiction- or industry-specific norms.
- **Uploads are not quote-verified.** Verification runs only on pasted text.
- **Five contract types.** Lease, NDA, employment, freelance, terms of service. Others produce lower-quality results.
- **English only.** Not validated on non-English contracts.

The bottom line

Fine Print is engineered to be reliable at what it is: a fast, transparent, uncertainty-aware triage aid. It will help a reader work through a contract more carefully and know what to question. It will not catch everything, and for any high-stakes agreement it is a complement to a qualified attorney, never a replacement. Where it is uncertain, it says so. That honesty is the feature.